

FarmNet: Identifying farms from satellite imagery.

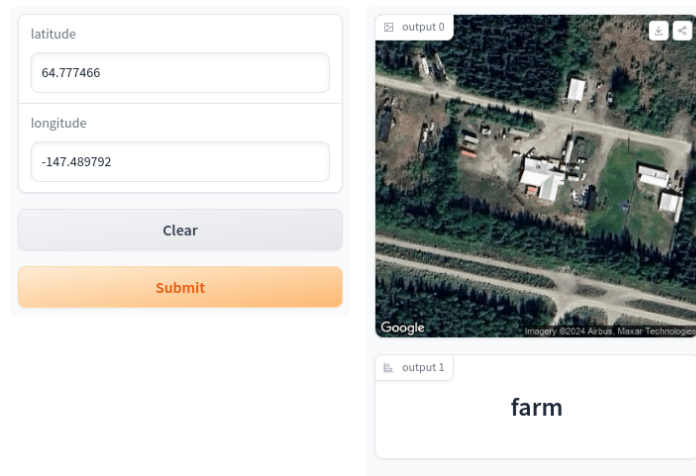
Filip Zawadka Group 51

Abstract

Due to the rise of automation and wealth inequality, I see a rise in factory farms, which contribute a significant public health, especially in neighboring communities. This project attempts to monitor the factory farms using satellite imagery.

1 Website

Making use of a simple CNN and Google Maps Api I created a [website](#) to identify factory farms. after imputing longitude and latitude of wanted map point, image of neighboring area will be displayed and identified whether it's a factory farm or not.



The screenshot displays the FarmNet web application interface. On the left, there is a form with two input fields: 'latitude' containing the value '64.777466' and 'longitude' containing '-147.489792'. Below these fields are two buttons: a grey 'Clear' button and an orange 'Submit' button. To the right of the form, there is a satellite image labeled 'output 0' showing a farm with several white buildings and green fields. Below the image, there is a white box labeled 'output 1' containing the word 'farm' in bold black text. The image has a 'Google' logo and 'Imagery ©2024 Airbus, Maxar Technologies' text at the bottom.

Figure 1: Recognized farm.

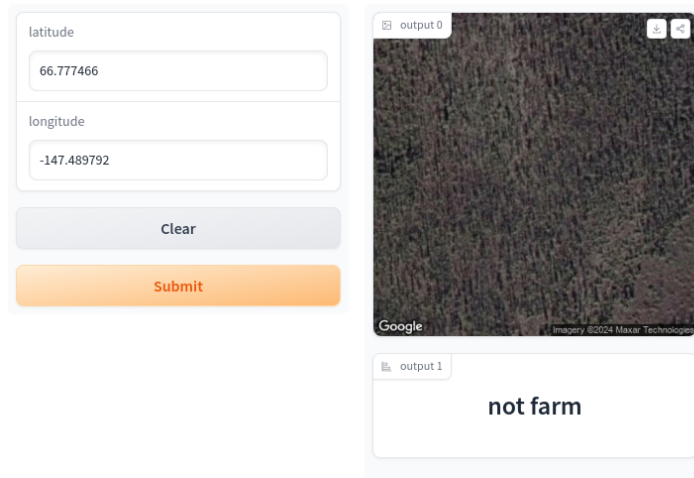


Figure 2: Forrest recognized as not a farm

2 Data

To generate data I use selenium to scrap data points from [counterglow](#), it's an amazing project where factory farms in the United States are submitted and later verified. I need to iterate over those points to collect latitudes and longitudes of marked points and corresponding labels, which include: "Slaughterhouse", "Broiler (Meat) Chickens", "Layer Hens (Eggs)", "Cows (Dairy, Meat)", "Cows (Meat)", "Pigs (Meat)", which provides us with more information about the farm. Points are constantly updated, so to improve the dataset, they should be collected periodically.

After collecting the points I use Google Maps Static Api, using satellite imagery to access the surrounding area of a data point at zoom "17", which should provide the most information about the area, while including enough of detail.

After this process I should have the "true" labels, to generate "false" labels I generate random points and add them to the dataset. I use 3:1 ratio of random points to farms, so it's not overwhelmingly beneficial to choose one label over the other.

3 Model Architecture

I started with resnet18, and fine-tuned it on our data, but very quickly observed that the model over-fitted to the data, seeing validation accuracy getting much lower than the training accuracy. After I expand the dataset, I might come back to this architecture, but for now I decided to use a simpler architecture.

3 Layers of 2d convolutional layers followed by max pooling and then by 2

linear layers, were used, so we have a simple enough architecture, for the amount of data. I also limited the classification to a binary one, to simplify the task for the model. This led to 0.95 accuracy on the test set.

4 How to run

To run the whole pipeline, first, use the `farms_scraper.ipynb`, it will collect points from counterglow and append them to `farms_scraped.csv` file, which all current datapoints

Next, `generate_satellite_images_google_api.ipynb`, this will make use of the scraped datapoints to generate google maps satellite images, it will also add random datapoint images.

Finally, to train the model use `train_cnn.ipynb`, which will train the model, and add it to wandb model registry. The production tag is set to "latest", so the model should be used on the website, as soon as it reboots.